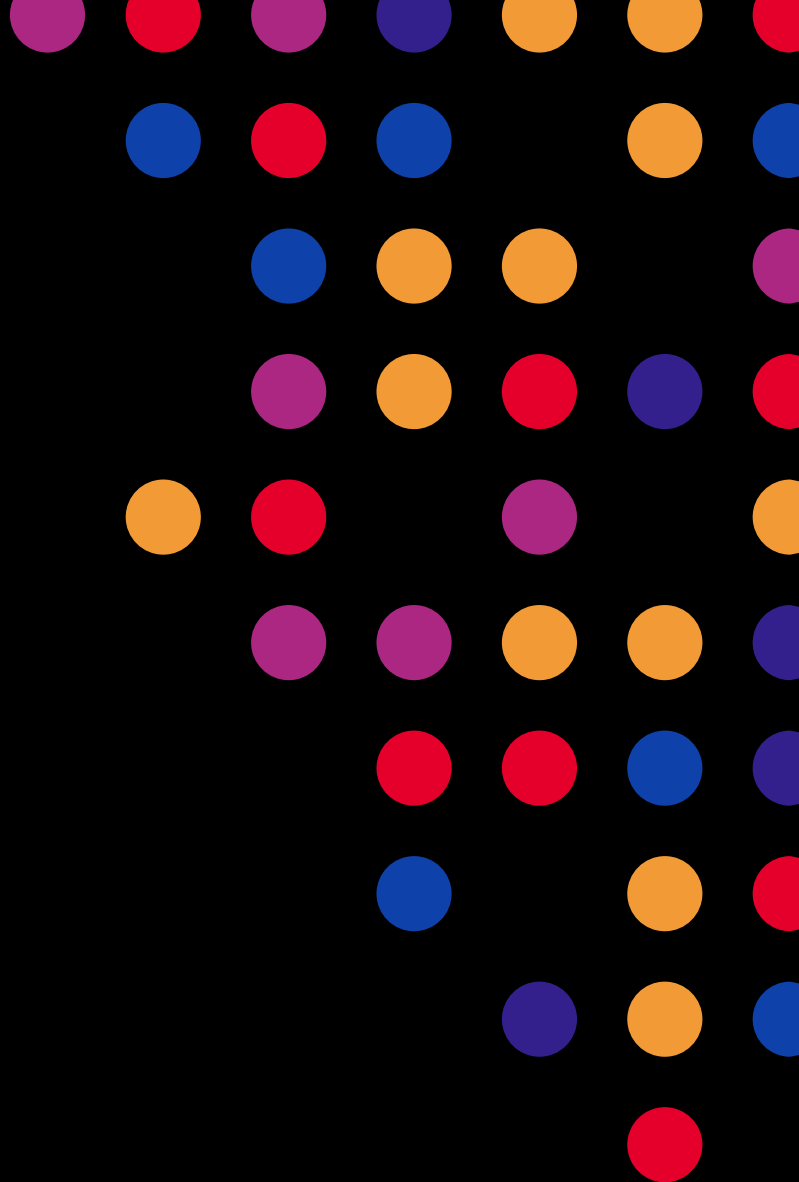




The Anatomy of a Distributed Trace

Dave McAllister

Senior Open Source Technologist
NGINX

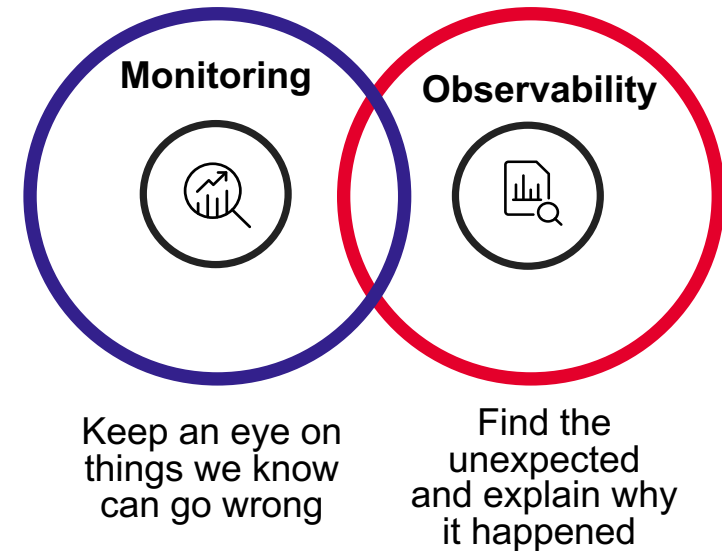
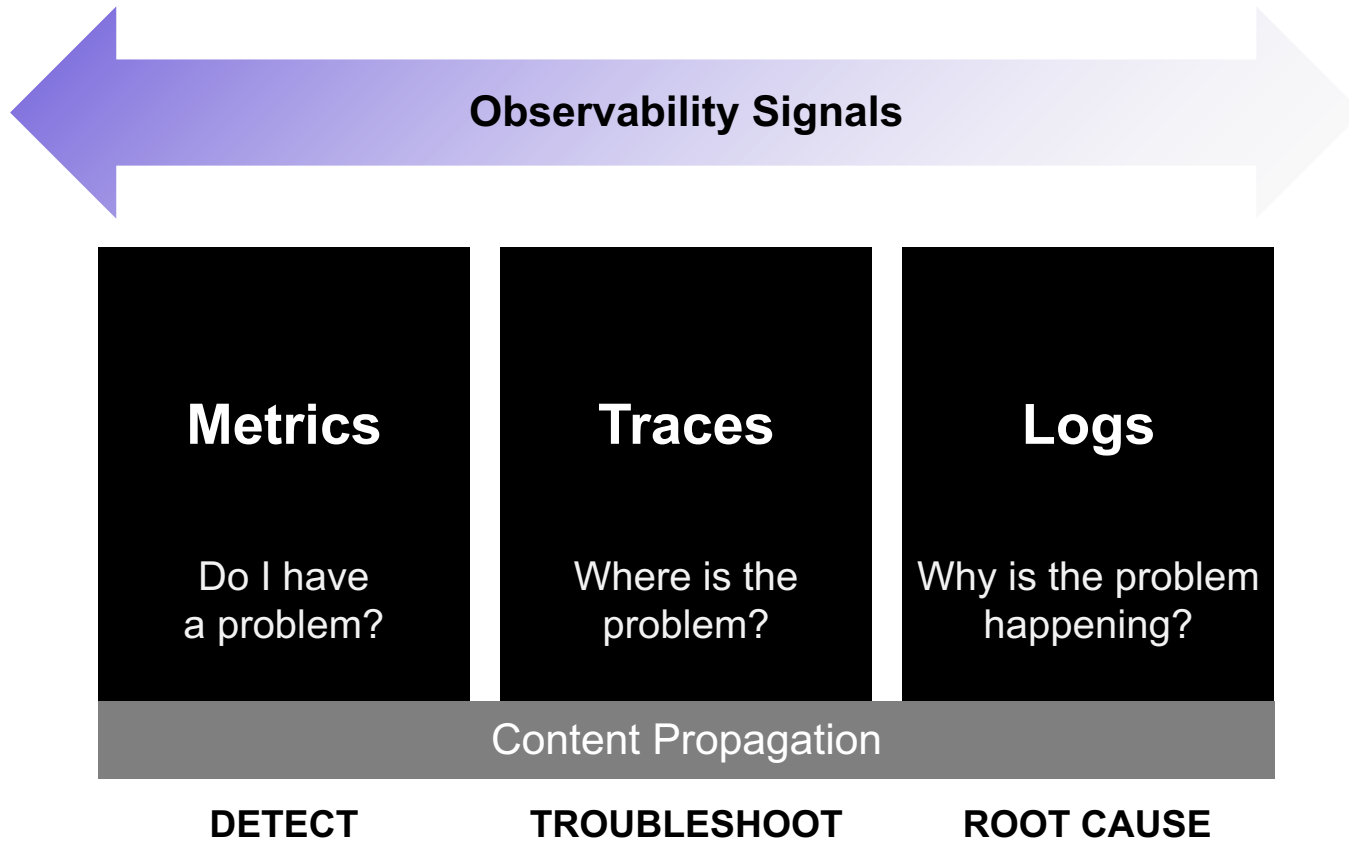


Observability is a data problem

The more observable a system, the quicker we can understand why it's acting up and fix it.

Observability

Observability helps detect, investigate and resolve the "unknown unknowns"—fast



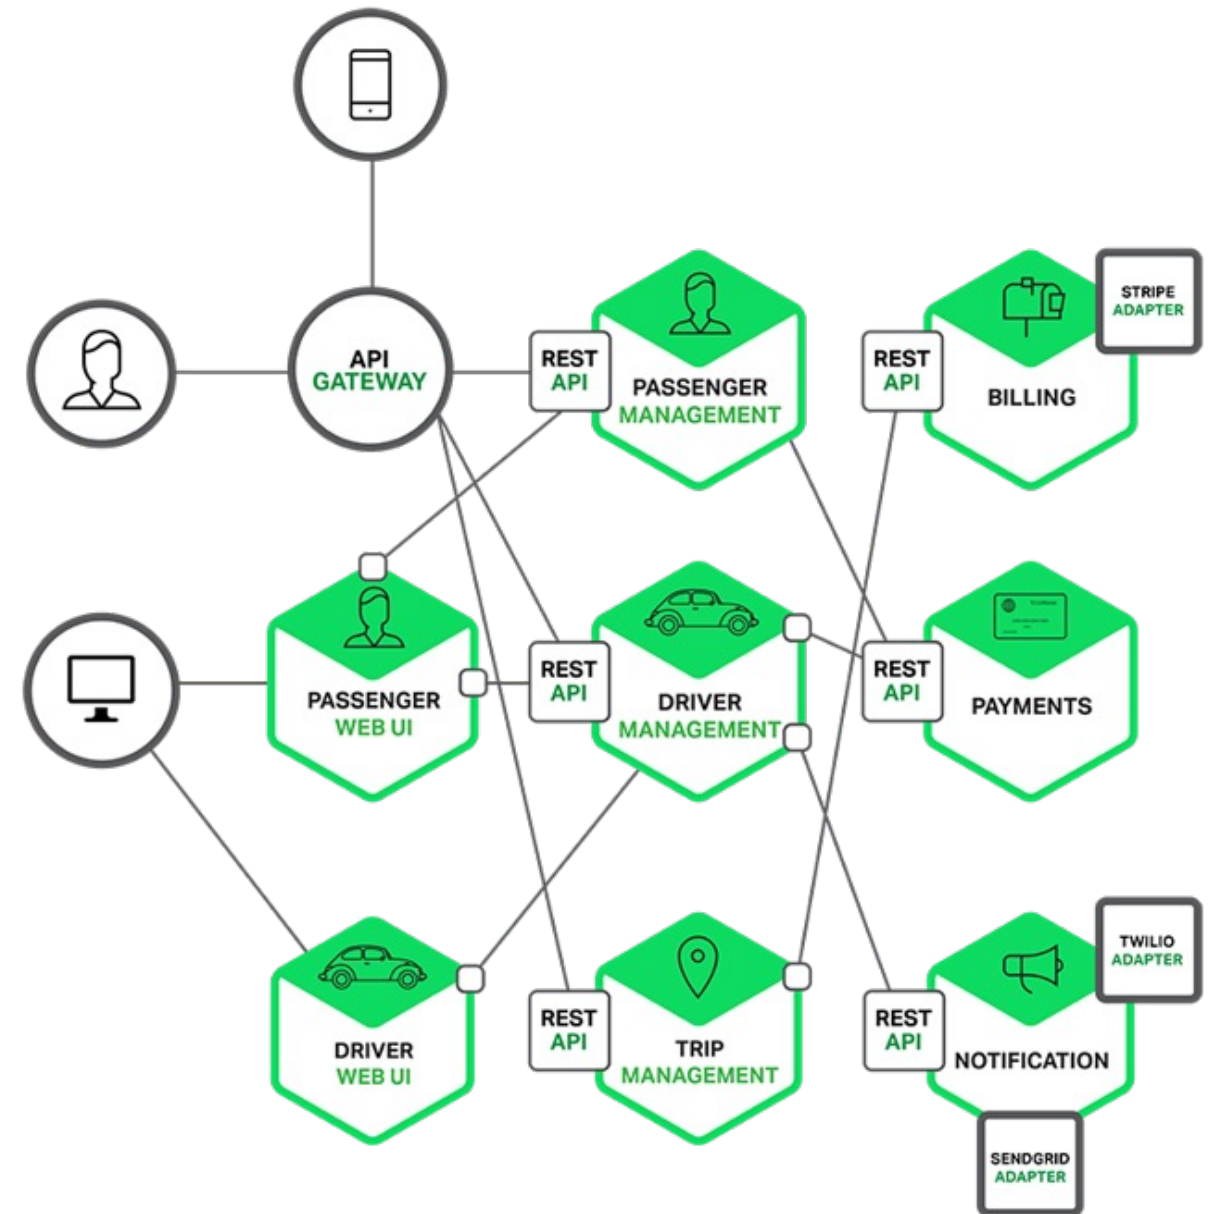
- Better visibility to the state of the system
- Precise and predictive alerting
- Reduces mean time to clue (MTTC) and mean time to resolution (MTTR)

So why observability?

Microservices!

A single application composed of many loosely coupled and independently deployable smaller services

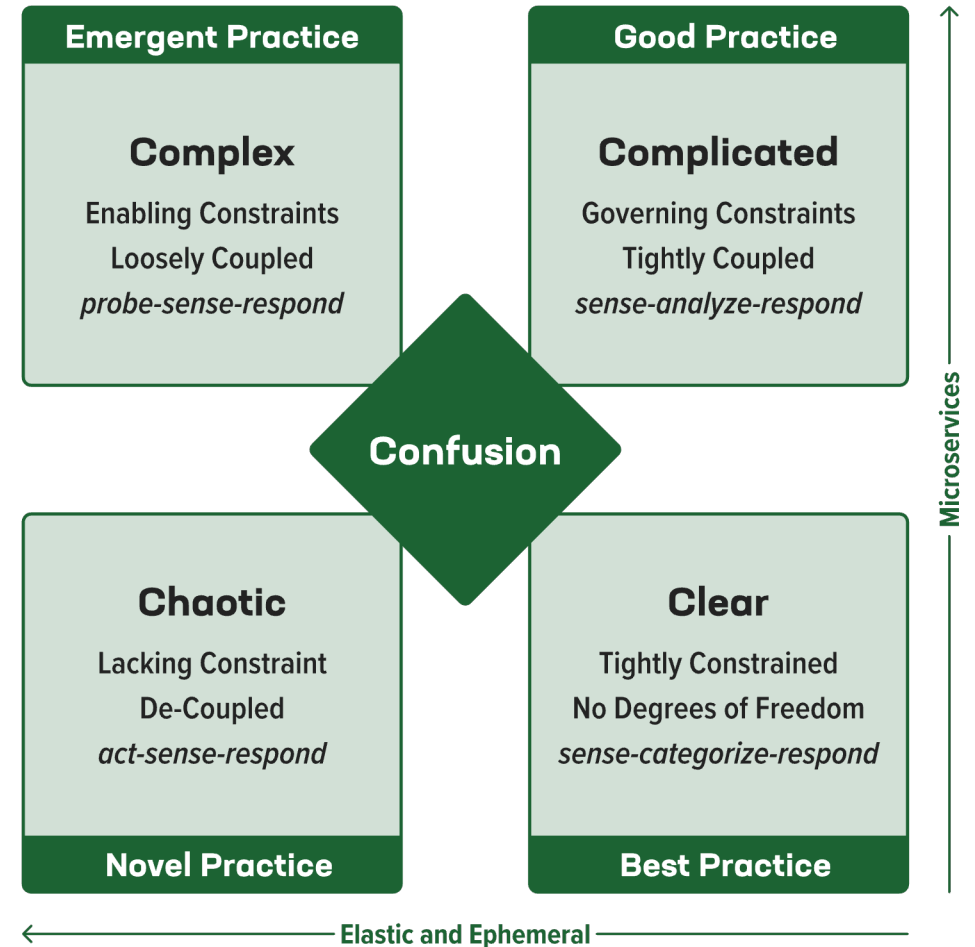
- Often polyglot in nature
- Highly maintainable and testable
- Loosely coupled
- Independently deployable
- Often in cloud environments
- Organized around capabilities



Microservices add challenges

Especially in a cloud environment

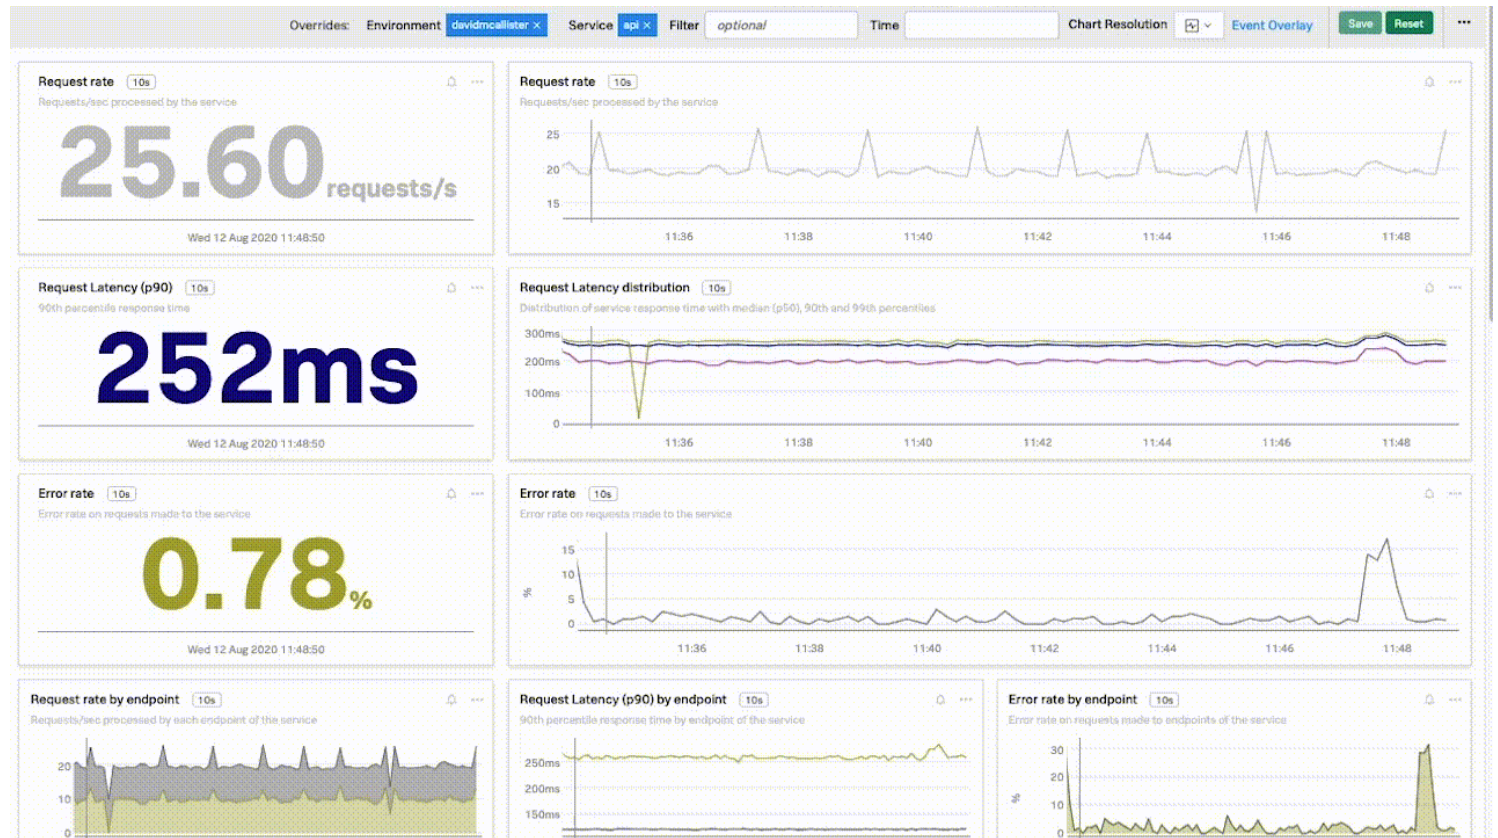
- Microservices create complex interactions
- Failures don't exactly repeat
- Debugging multitenancy is painful
- So much data!



Tracing is a data problem

What's distributed tracing good for?

- Tracks requests
- Provides actionable insights into app/user experiences
- Defines additional metrics for alerting, debugging
- Rapid MTTC, MTTR



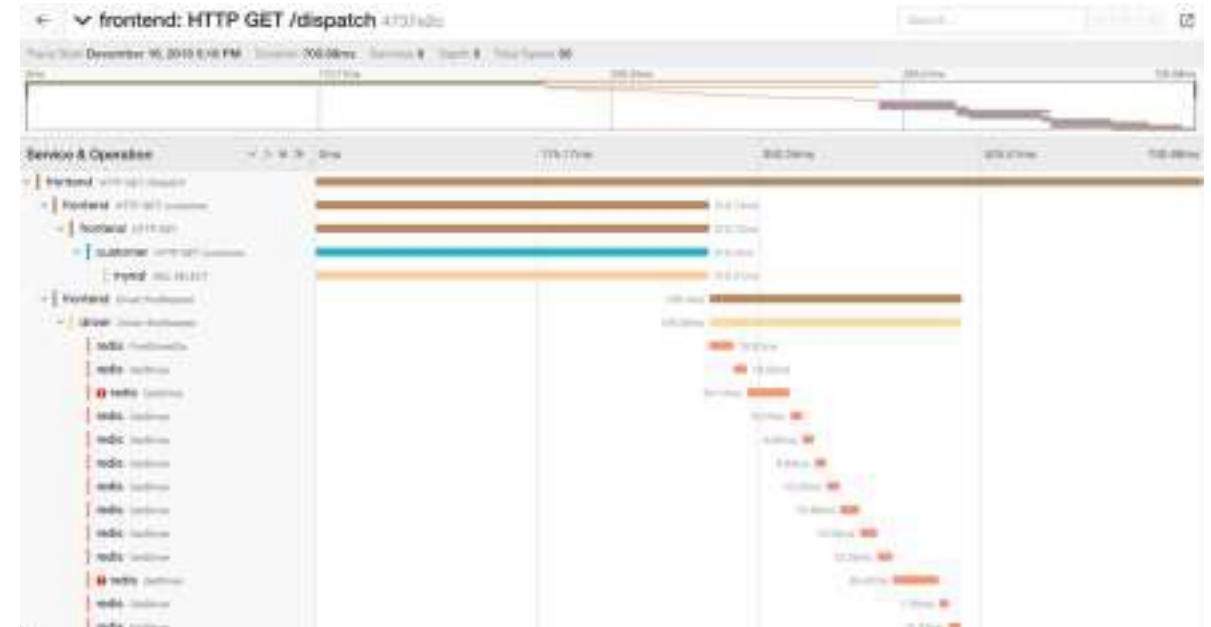
RUM, Synthetics, NPM, APM, infrastructure

RUM, Synthetics, NPM, APM, infrastructure

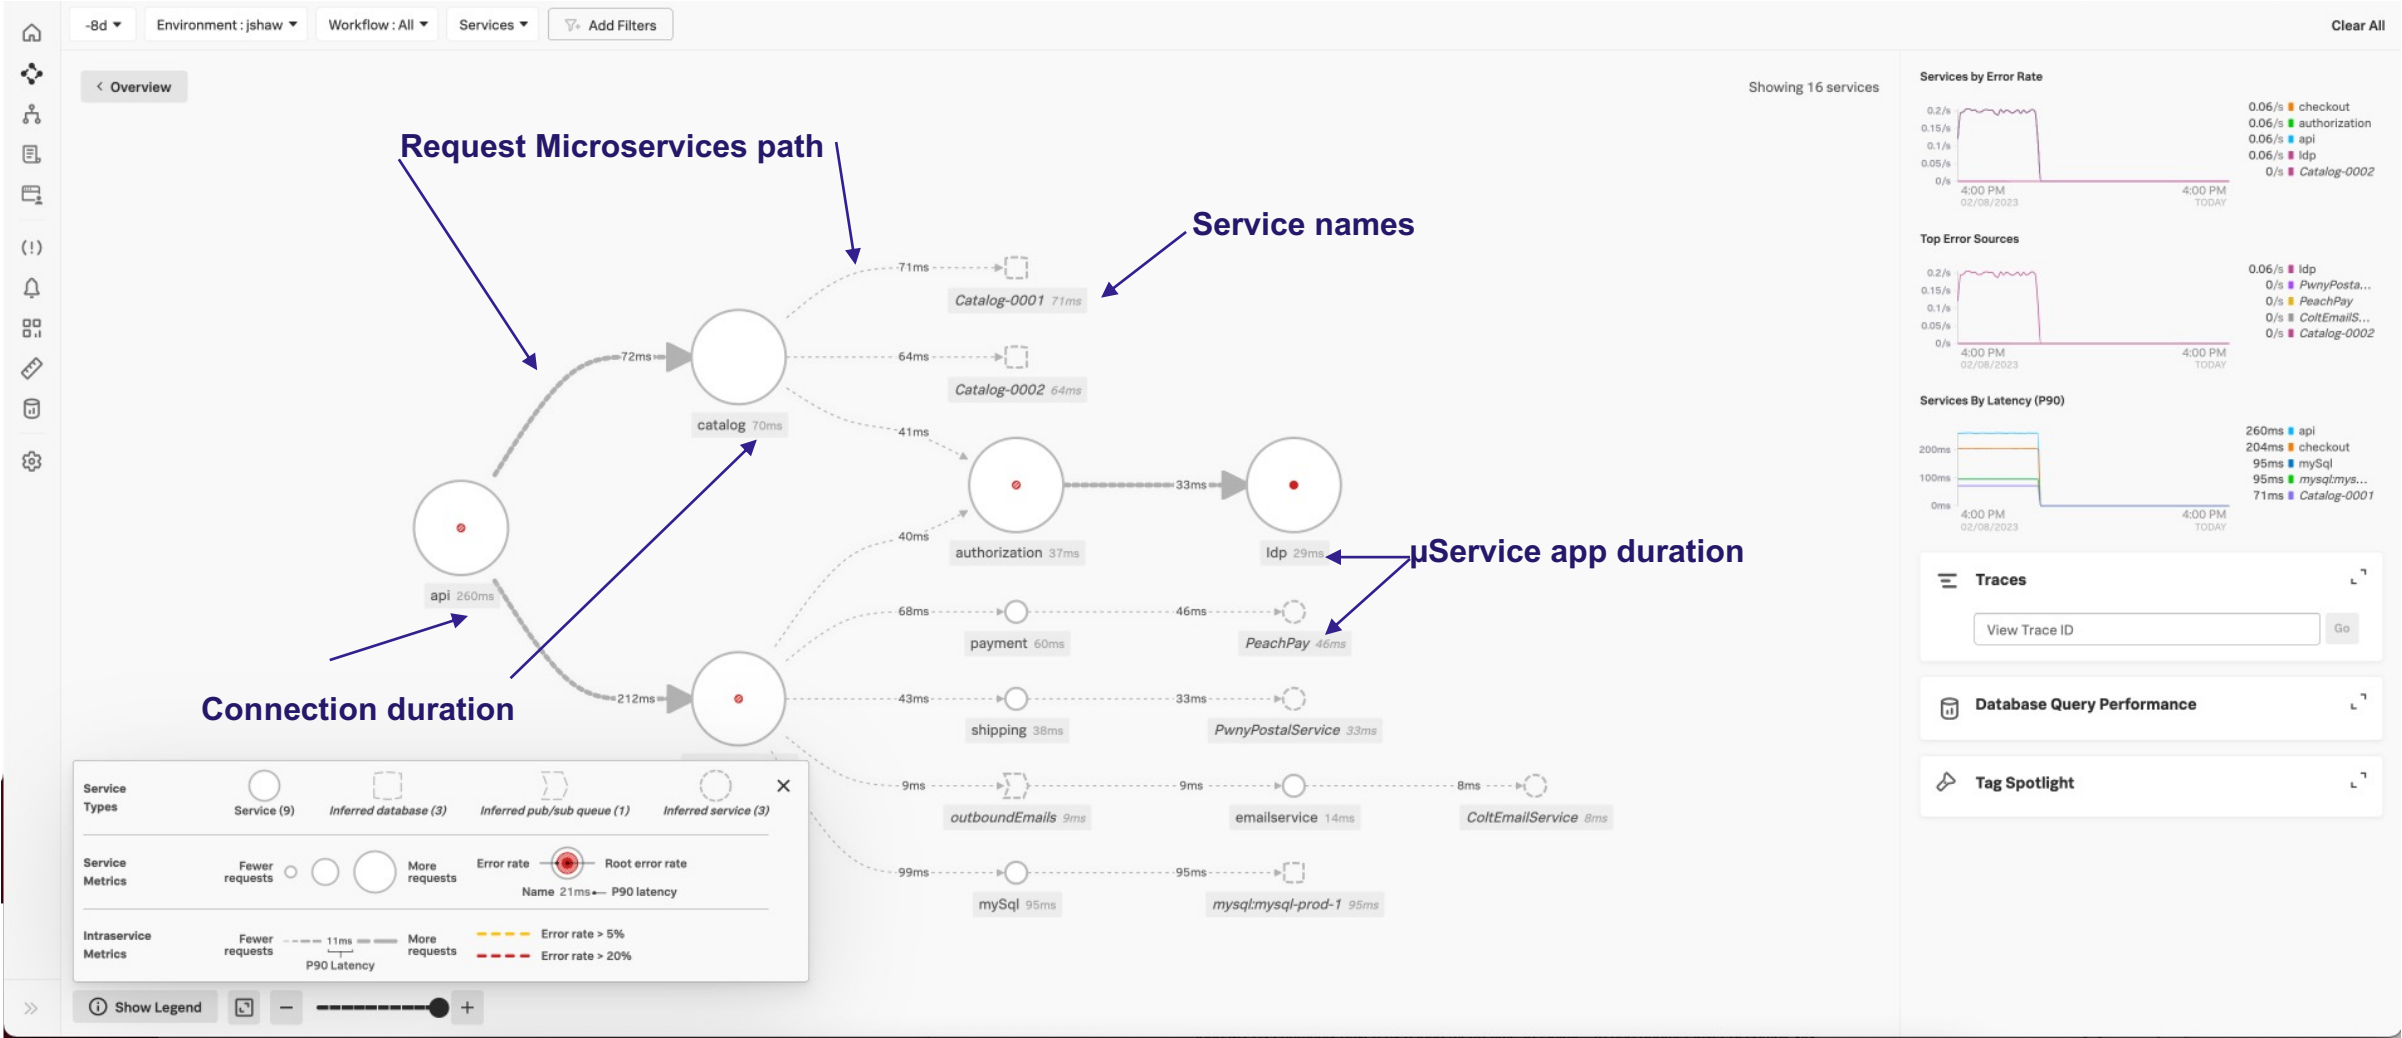
© 2024 F5

Tracing concepts


- Span
 - Represents a single unit of work in a system
- Trace
 - Defined implicitly by its spans
- Distributed context
 - Tracing identifiers
 - Tags
 - Options that are propagated from parent to child spans



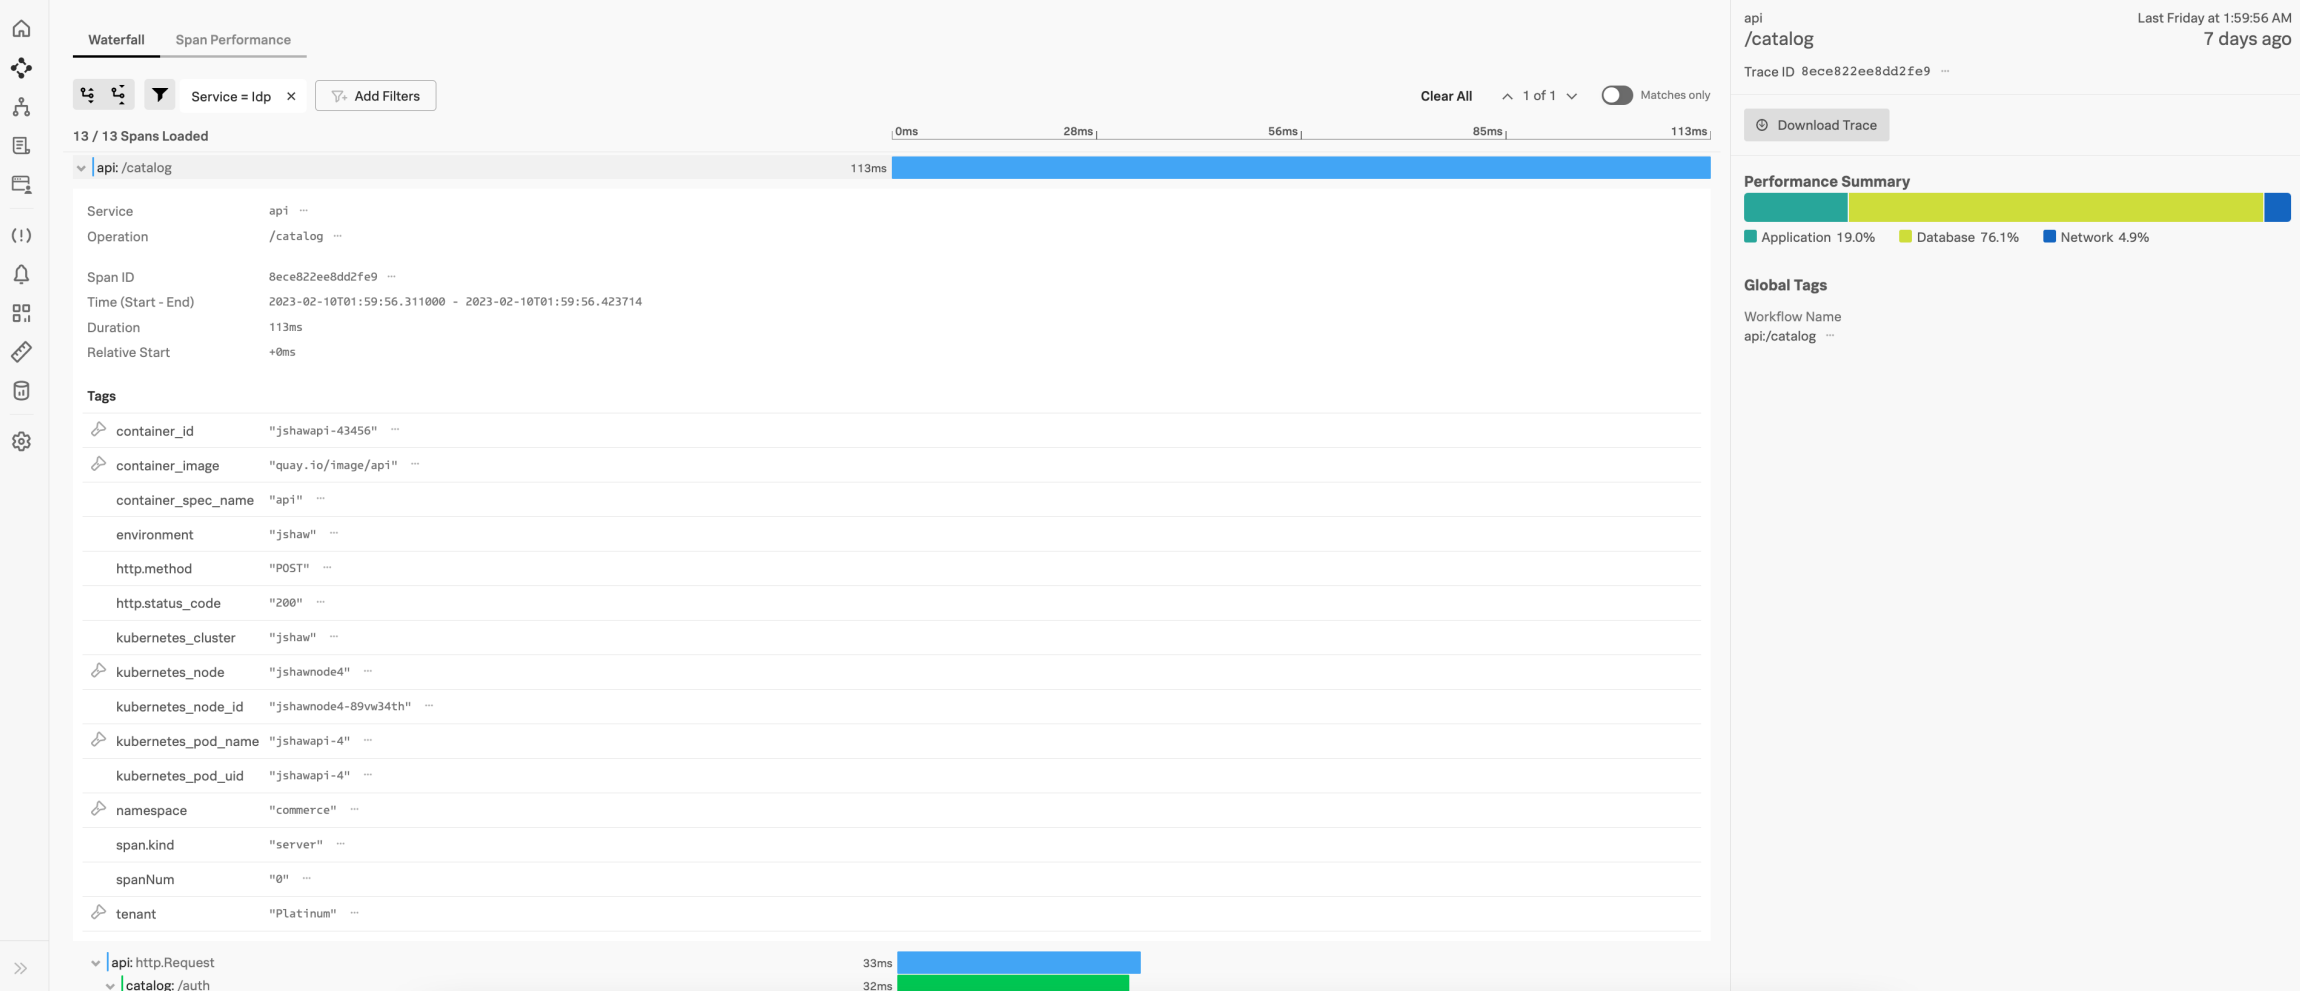
Let's look at a trace




A different way of looking at a trace



Observability includes baggage



2024-01-03 19:06:46.294 PUT 202 /api/v1/orgs/21d88bc4-5302-452d-ba42-492cee33b341/groups/RG-DRK-NGX-SHD-EUS-01/deployments/NGX-DRK-NGX-SHD-EUS-01/config



Service & Operation	0µs	2.07s	4.14s	6.21s	8.28s
ARP /api/v1/orgs/:org_id/groups/:group_name/deployments/:deployment_name/config (191.41ms)	191.41ms				
Middleware (191.06ms)	191.06ms				
AnalyzeNginxConfig.Do (26.14ms)	26.14ms				
ngxDeployments.find (17.49ms)	17.49ms				
ngxStatuses.find (4.69ms)	4.69ms				
HTTP blob PUT (15.34ms)	15.34ms				
HTTP blob HEAD (9.61ms)	9.61ms				
ngxDeployments.update (86.31ms)	86.31ms				
ngxStatuses.find (4.68ms)	4.68ms				
ngxStatuses.update (15.1ms)	15.1ms				
DeploymentDesiredStateSender.Send (442.27ms)	442.27ms				
send feed.deployments.desired_state (365.28ms)	365.28ms				
MTR receive (30.44ms)	30.44ms				
handle feed.deployments.desired_state (30.33ms)	30.33ms				
DPO receive (334.21µs)	334.21µs				
handle feed.deployments.desired_state (5.39s)	5.39s				
operator.handleDesiredState (5.38s)	5.38s				
zones.eligibilityCalculator.GetEligibleZones (1.13s)	1.13s				
resourcegroup.Controller.syncResourceGroup (55.5µs)	55.5µs				
agentcert.Controller.syncAgentCert (38.5µs)	38.5µs				
frontend.Controller.update (4.18s)	4.18s				
operator.waitForControllers (4.18s)	4.18s				
deployment.syncInstanceGroups (1.02s)	1.02s				
autoscale.Controller.syncAutoScaleSettings (13.1µs)	13.1µs				
nginxconfig.Controller.syncConfig (1.02s)	1.02s				
send feed.deployments.status (19.7ms)	19.7ms				
ARP receive (186.44ms)	186.44ms				
handle feed.deployments.status (186.27ms)	186.27ms				
ngxDeployments.find (18.82ms)	18.82ms				
UpdateDeploymentFromStatus (132.71ms)	132.71ms				

Middleware (191.06ms)

AnalyzeNginxConfig.Do (26.14ms)

nginxDeployments.find (17.35ms)

Middleware

nginxDeployments.find

Service: ARP | Duration: 191.06ms | Start Time: 122.25µs (19:06:46.294) | Child Count: 9 | Kind: internal | Status: unset

Service: ARP | Duration: 17.35ms | Start Time: 1.28ms (19:06:46.296) | Kind: client | Status: unset

Logs for this span

Logs for this span

Span Attributes

Span Attributes

clientType "Uncategorized"

cluster "production-f5nalbeus01"

dnslabel "f5nalbeus01"

fleetenv "production"

location "eastus"

resourceID "/subscriptions/21d88bc4-5302-452d-ba42-492cee33b341/resourceGroups/RG-DRK-NGX-SHD-EUS-01/providers/NGINX.NGINXPLUS/nginxDeployments/NGX-DRK-NGX-SHD-EUS-01"

userAgent "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/120.0.0.0 Safari/537.36 azsdk-net-Microsoft.Liftr.Nginx.SaaS.Client/0.2102.02545-build731 (.NET 6.0.25; Linux 5.15.138.1-4.cm2 #1 SMP Thu Nov 30 21:48:10 UTC 2023)"

Resource Attributes: service.name = ARP | service.version = 1.20231212.1104728472 | telemetry.sdk.language = go | telemetry.sdk.name = opentelemetry...

Events (1)

148.25µs: message = user data added to context

Log timestamps are relative to the start time of the full trace.

SpanID: cd3f573eb6e6f62a



What do we need for tracing?

- Generated unique IDs
- Auto-propagation
- Telemetry consolidation
- Distributed environment capable
- Standards-based agents, cloud integration
- Automated code instrumentation
- Support for developer frameworks
- Any code, any time



Tracing API concepts – OpenTelemetry

TracerProvider is the entry point of the API. It provides access to Tracers.

- Stateful object holding configuration with a global provider and possibly additional ones.

Tracer is the class responsible for creating Spans.

- Named and optionally versioned with each instrumentation library using values guaranteed to be globally unique.
- Delegates getting active Span and marking a given Span as active to the **Context**.

Span is the API to trace an operation.

- Immutable **SpanContext** represents the serialized and propagated portion of a Span.

Enabling distributed tracing

Two basic options:

- Traffic inspection (e.g., service mesh with context propagation)
- Code instrumentation with context propagation

Focusing on code:

- Add a client library dependency
- Focus on instrumenting all service-to-service communication
- Enhance spans (key-value pairs, logs)
- Add additional instrumentation (integrations, function-level, asynchronous calls)

What does this enabling mean?

Traces

1. Instantiate a tracer
2. Create spans
3. Enhance spans
4. Configure SDK

Automatic Instrumentation

- Just add the appropriate files to the app.
This is language dependent

Manual Instrumentation

- Import the OTel API and SDK
- Configure the API
- Configure the SDK
- Create your traces
- Create your metrics
- Export your data

What problems does tracing clarify?

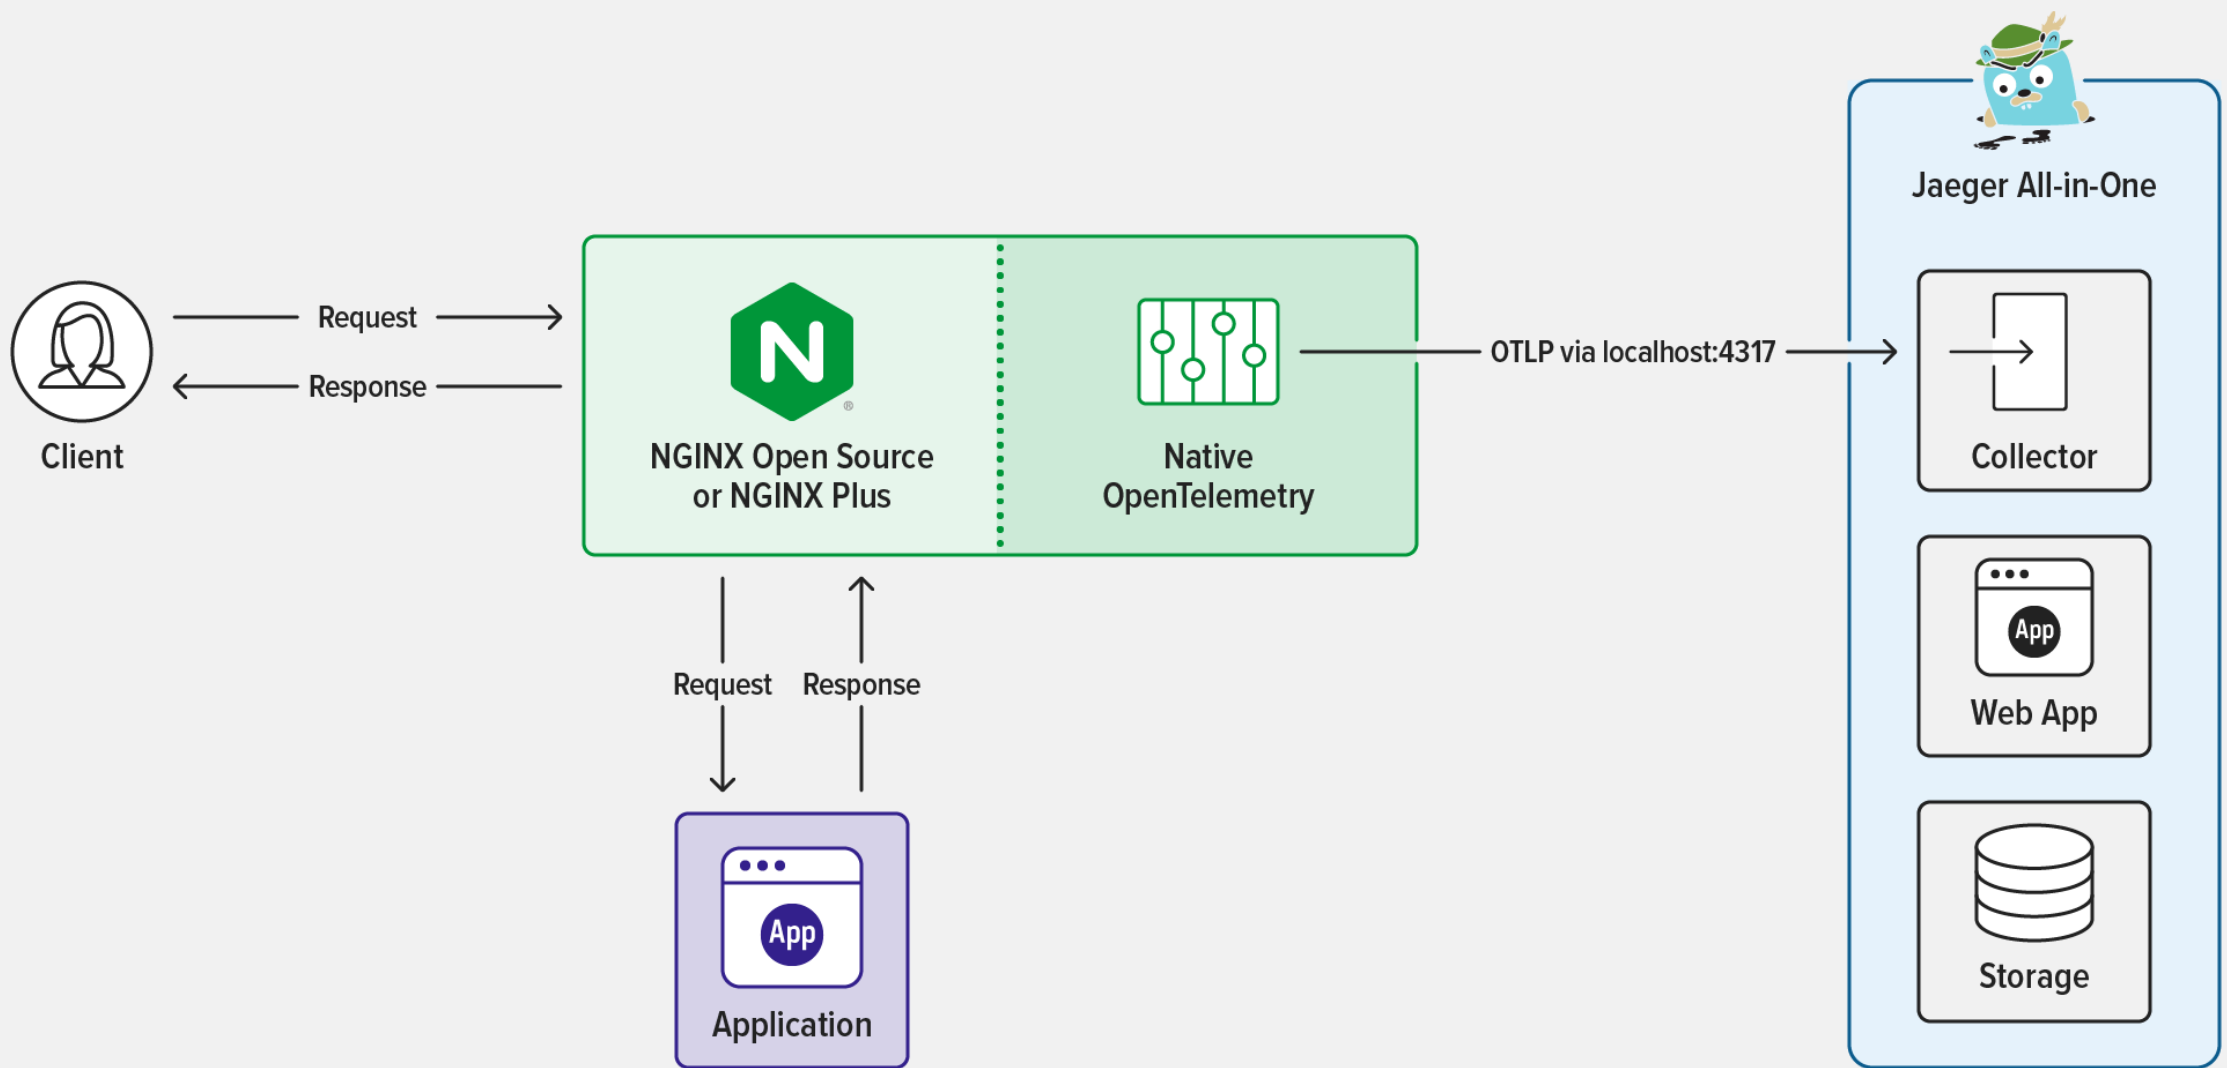
A tiny subset of answers

- Performance issues
- Mean time to resolution and detection
- Context for metrics and logs
- Connections between infrastructure and applications

So Why a Native OpenTelemetry Module for NGINX?

[nginxinc/nginx-otel \(github.com\)](https://github.com/nginxinc/nginx-otel)

- Enables NGINX to send to an Otel collector
 - Fully supports W3C trace context
 - Supports OTLP and gRPC protocols
- Performance
 - Current community modules reduce request processing by ~50%
 - Native module ~10%
- Setup and config are inline with the NGINX application configuration
- Allows dynamic control of trace parameters using cookies, tokens and/or variables
- Prebuilt packages are available, including RedHat and derivatives, Debian, Ubuntu and derivatives



```

load_module modules/nginx_otel_module.so;

events {
}

http {

    otel_exporter {
        endpoint localhost:4317;
    }

    server {
        listen 127.0.0.1:8080;

        location / {
            otel_trace          on;
            otel_trace_context inject;

            proxy_pass http://backend;
        }
    }
}

```

Example Configuration

- Otel_exporter – specifies Otel data export parameters
 - Endpoint – address to accept telemetry data
 - Interval – interval (max) between exports (5s)
 - Batch_size – max spans sent in one batch per worker (512)
 - Batch_count – number of pending batches per worker (4)
- Otel_trace – enables or disables tracing
 - Can be a variable
- Otel_trace_context – propagation directives
 - Extract | inject | propagate | ignore

Automatically add span attributes

- http.method
- http.target
- http.route
- http.scheme
- http.flavor
- http.user_agent
- http.request_content_length
- http.response_content_length
- http.status_code
- net.host.name
- net.host.port
- net.sock.peer.addr
- net.sock.peer.port

Simple tracing example – all HTTP requests

```
http {  
    otel_exporter {  
        endpoint localhost:4317;  
    }  
  
    otel_trace on;  
  
    server {  
        location / {  
            proxy_pass http://backend;  
        }  
    }  
}
```


Parent-based tracing – inherit and record only if parent is sampled

```
http {  
    server {  
        location / {  
            otel_trace $otel_parent_sampled;  
            otel_trace_context propagate;  
  
            proxy_pass http://backend;  
        }  
    }  
}
```

Ratio-based – example of percent of traffic

```
http {
    # trace 10% of requests
    split_clients $otel_trace_id $ratio_sampler {
        10%      on;
        *        off;
    }

    # or we can trace 10% of user sessions
    split_clients $cookie_sessionid $session_sampler {
        10%      on;
        *        off;
    }

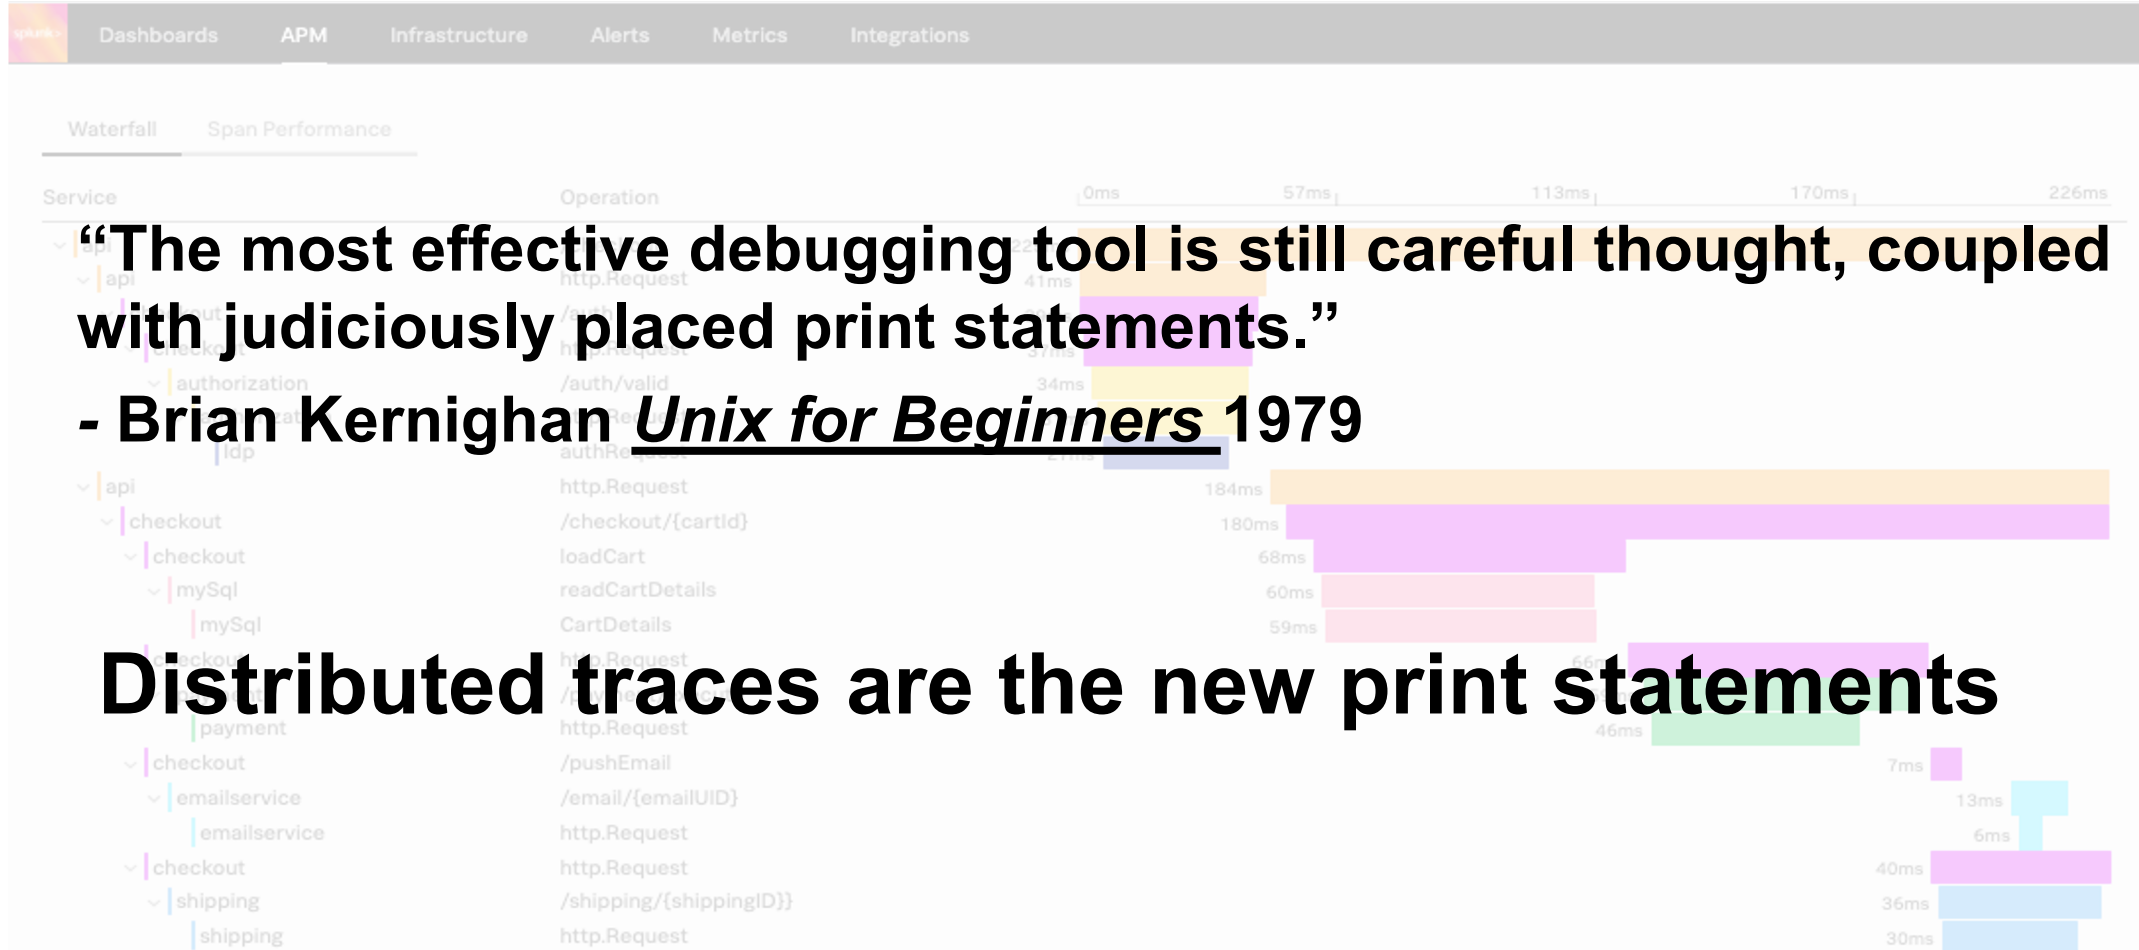
    server {
        location / {
            otel_trace $ratio_sampler;
            otel_trace_context inject;

            proxy_pass http://backend;
        }
    }
}
```

Distributed tracing summary


- Gives insights into the app and its infrastructure
- Requires effort to return value
 - The “win” is when the project is already OTel instrumented
- Is a good “user happiness” proxy
- Does not magically solve your issues

Closing thought



Odysseus is in orbit, will attempt to touch down at 5:30pm
Eastern time

Thanks!



[LinkedIn: in/davemc](https://www.linkedin.com/in/davemc)

[Slides on GitHub](#)

